

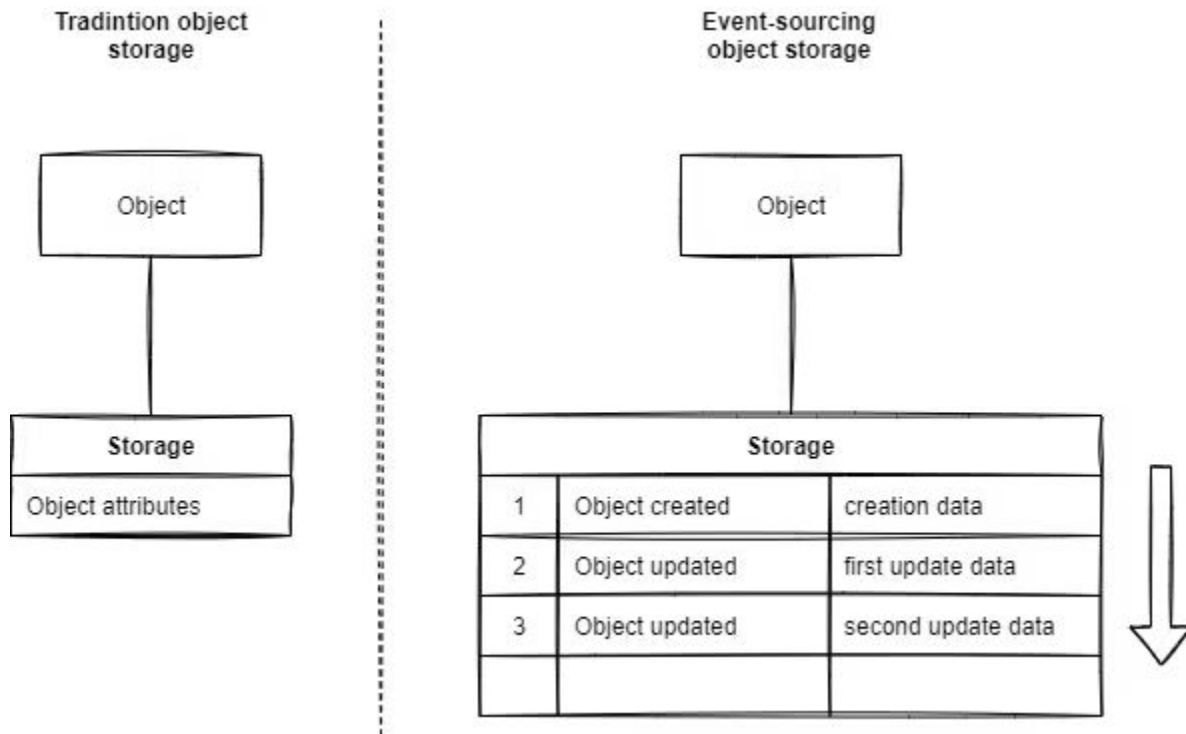
Event Sourcing using Akka Persistence

Agenda


- What is Event Sourcing (ES)?
- What is CQRS?
- Features of ES
- Challenges with ES
- ES with Akka Persistence

What is Event Sourcing?

- ★ ES is a design pattern for storing the **internal state** of a system as a **sequence of events**
- ★ In ES we don't persist the current state, but **changes** captured in events that lead to this state



- ★ **The state** of the system is derived by **replaying these events** in order in which these events were produced



ES Journal

- ★ Journal is a sequence of events, known as event log or event store
- ★ Each event is identified by a unique ID, sequence number and it represents a change that was made to the system state over time
- ★ Journal events are always **immutable and append-only**
- ★ Journal is the **source of truth**

Commands vs Events vs Aggregate

- ★ **Command** is a **change request** to the system
- ★ **Event** is the actual **fact** that has happened in the past
- ★ An **aggregate** is a collection of related domain objects
 - Used to enforce consistency and integrity within the system
 - Used to handle commands and generate events
- ★ When a new command is handled, the system can produce
 - ... one or more events
 - ... or an error

When a command is received, the system validates the command, enforcing business and logical constraints to ensure that only **valid and authorized** events are generated and written to the journal

Current application state

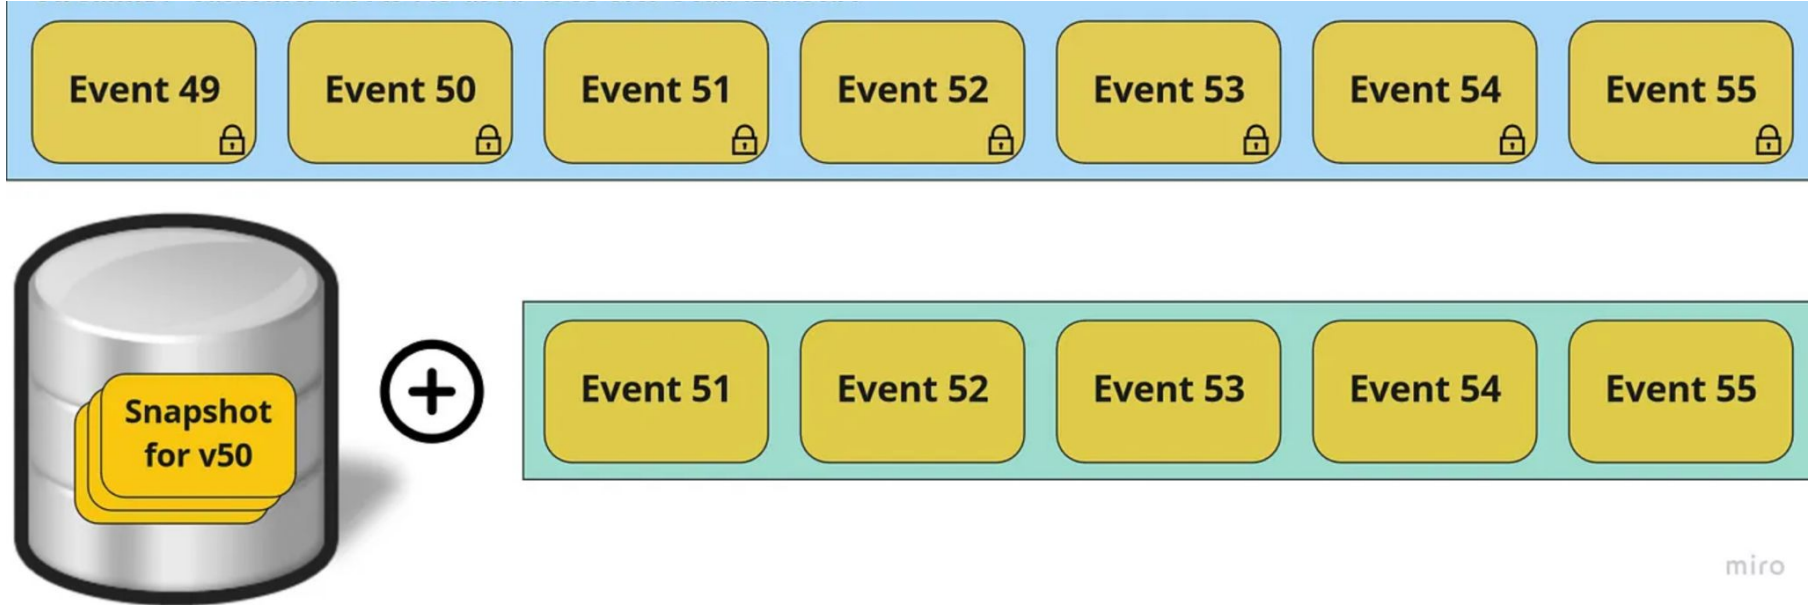
- ★ Is built based on event journal
- ★ can be discarded completely and rebuilt at any time
- ★ can be used to validate commands prior generating events because events should not lead to inconsistent state

Example of “Bank Account” aggregate

| Command in order | Event(s) or Error(s) | State |
|---------------------|--|----------------|
| 1: Increase on 100 | Event: Increased on 100 | Balance is 100 |
| 2: Decrease by 30 | Event: Decreased on 30 | Balance is 70 |
| 3: Decrease by 100 | Error: Balance falls below 0 | unchanged |
| 4: Increase on 2000 | Error: Balance limit of 1000 is reached | unchanged |

Snapshots

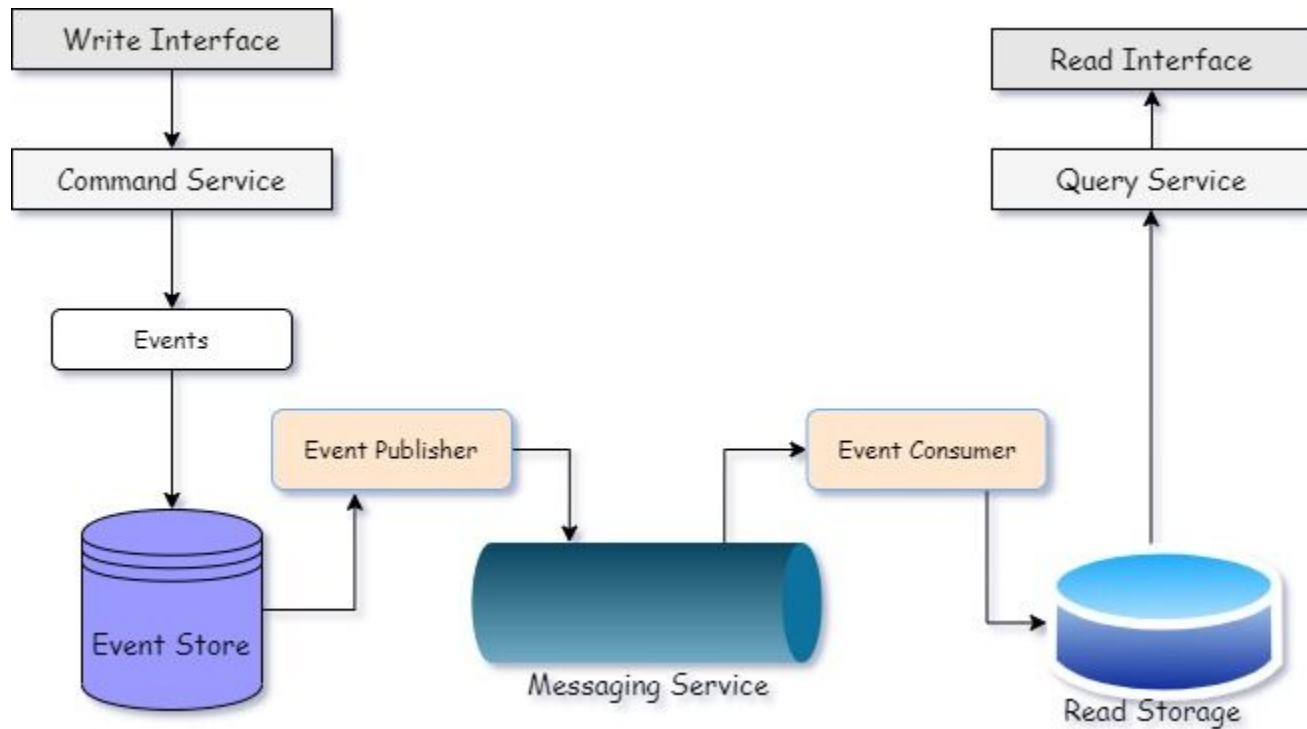
- ★ If there're a lot of events, replaying events to build state can be slow
- ★ Creating snapshots each N-events is a nice optimization
- ★ In case of using snapshots, current state is calculated by
 - Restoring the state to the latest snapshot
 - And then replaying remaining events, that were generated after the snapshot creation, on top



miro

What is CQRS?

- ★ Command Query Responsibility Segregation
- ★ CQRS is a design pattern that separates handling read and write operations
- ★ CQRS separates reads and writes into different models, using **commands** to update data, and **queries** to read data



Write Model

- ★ Is responsible for handling commands to the system
- ★ Command and Aggregate Handlers:
 - the command handler is responsible for receiving commands and passing them to the appropriate aggregate
 - the aggregate handler is responsible for processing commands and generating events on valid commands

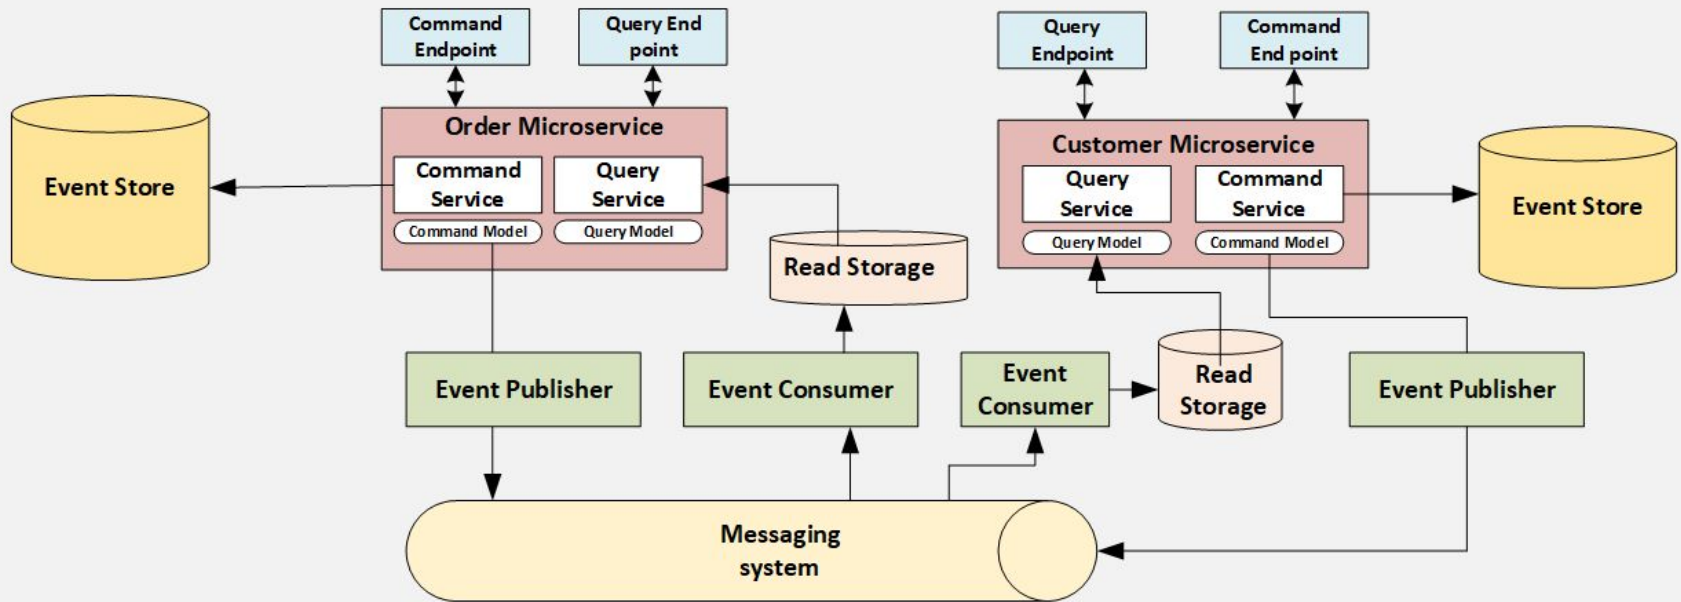
Read Model

- ★ Is responsible for handling read requests to the system and generating views or projections
- ★ Is optimized for fast queries and can be updated asynchronously and eventually in the background

Views

- ★ A view is a representation of the state that is optimized for specific use cases
- ★ Views can be stored as materialized views
- ★ A materialized view is a precomputed view that is stored in a separate datastore for faster access
- ★ Materialized views can be updated in real-time or on a scheduled basis

Microservices with CQRS and Event Sourcing



Features of ES

- ★ Rebuild state at any point of time
- ★ Avoids relational structure of the data
- ★ Compatibility with CQRS & Domain Driven Design


Immutable event log

- ★ provides a complete history of all changes to the domain model over time, easy to trace history of changes (who did, what and when)
- ★ easy to build views for audit log and demonstrate compliance with regulatory requirements

Challenges of ES

- ★ Increased development and operational complexity
- ★ Evolution of system, events and data models
- ★ Choice of an aggregate root

Operational and development overhead



When your database admin
hears about event sourcing

Operational and development overhead

- ★ These approaches require additional infrastructure to handle
 - event stores for a journal
 - message brokers
 - materialized views
- ★ This can require additional resources and management overhead to maintain

Operational and development overhead

Many tips to optimize solution

- Build snapshots of the state
- Configure TTL for operational data (events, snapshots)
- Move historical data to cheaper storages with cheaper disks
- Use cloud-based infrastructure
- Use event sourcing frameworks

Evolution of events and data models

- ★ Systems evolve over time without having to change the underlying data model, but event and snapshot formats can change, need to understand how to
 - Choice of data formats of events
 - deal with event evolution
 - and how long you keep backwards compatibility
 - Personal comment: exposing journal events is a bad idea

Choice of aggregate root

- ★ An aggregate is a collection of related domain objects that are used to enforce **consistency** and **integrity** within the system
- ★ **Too small** aggregate lead to need for cross-aggregate transactions
- ★ **Too large** aggregate lead to large state, complicated events, system **scales** only to some **limit**

- *Why did the event sourcing engineer refuse to go to the party?*
- *Because he was busy replaying all the events of the day!*
@ chat gpt

